

David Kuster
dave@talldave.net
@ofTallDaveFame

Hibernate Metrics

Paying Attention to the ORM
Behind the Curtain

About Me

- Grails dev ~4 years
- Java/JEE dev ~10 years prior
- Bootstrapped startups to Fortune 500
- Developer of Grails Command Center
- “You're really good at writing your codes.”
- my wife

Agenda

- The landscape
- The solutions
- The problems that the solutions have uncovered
- And some fixes

A Caveat

- Hibernate is a big topic
- I came to Hibernate through GORM
- There are definitely gaps in my knowledge on this topic

What GORM Promises



The Reality



And...

“Object-relational mapping is the Vietnam of Computer Science.”

- Ted Neward

I love the smell of object-relational impedance mismatch in the morning



Best vs Worst Case Scenarios



The Point Being

PAY ATTENTION TO THE DATABASE

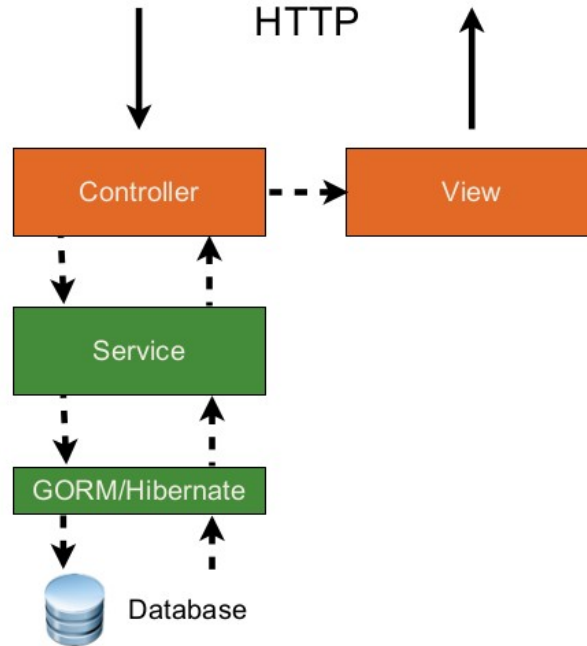
Any Sufficiently Advanced Technology Is Indistinguishable From Magic




HIBERNATE

COLLECTION NOT PROCESSED BY FLUSH? THROW SOME MORE EYE OF NEWT AT IT.

Conceptual Model



Database Access From Anywhere



Agenda

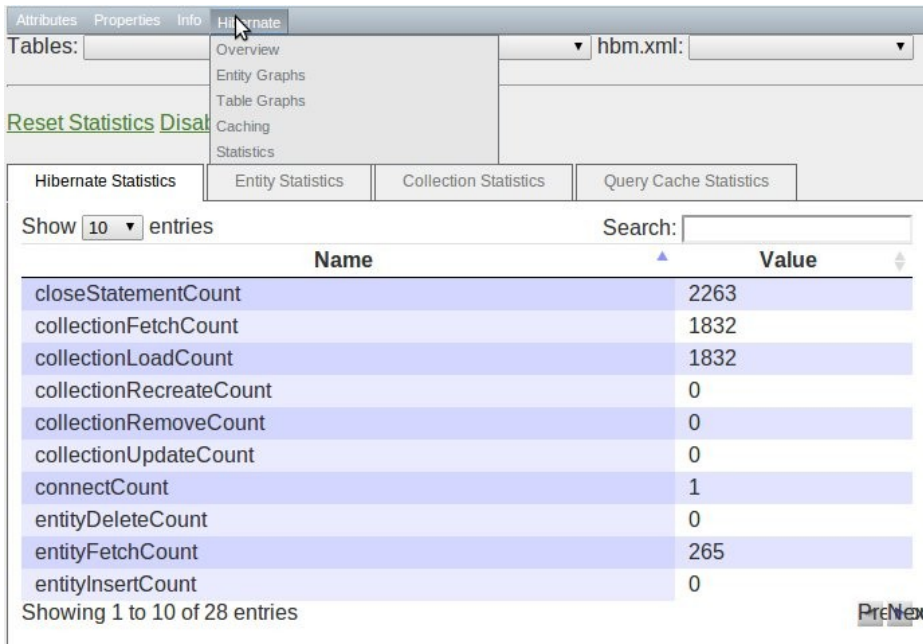
- The landscape
- **The solutions**
- The problems that the solutions have uncovered
- And some fixes

A Well-Trod Path

- Many have come this way before
- A number of plugins (and external/paid tools) that allow us to pay attention to the DB

Hibernate App-Info

- ❑ Extension to App-Info Plugin
- ❑ Add mixins to Config.groovy
- ❑ /appName/adminManage
- ❑ A ton of info, but not easily browsable



The screenshot shows the Hibernate App-Info web interface. At the top, there are tabs for 'Attributes', 'Properties', 'Info', and 'Hibernate'. The 'Hibernate' tab is selected, and a mouse cursor is pointing at it. Below the tabs, there is a 'Tables:' section with a dropdown menu showing 'hbm.xml:'. To the right of this, there is a 'Reset Statistics' link and a 'Disable' link. Below these links, there is a 'Hibernate Statistics' section with three sub-tabs: 'Entity Statistics', 'Collection Statistics', and 'Query Cache Statistics'. The 'Entity Statistics' sub-tab is selected. Below the sub-tabs, there is a 'Show' dropdown menu set to '10' and a 'Search:' input field. The main content area displays a table with two columns: 'Name' and 'Value'. The table contains 10 rows of statistics, with the first row highlighted in blue. The statistics are as follows:

Name	Value
closeStatementCount	2263
collectionFetchCount	1832
collectionLoadCount	1832
collectionRecreateCount	0
collectionRemoveCount	0
collectionUpdateCount	0
connectCount	1
entityDeleteCount	0
entityFetchCount	265
entityInsertCount	0

At the bottom of the table, it says 'Showing 1 to 10 of 28 entries'. In the bottom right corner, there is a 'PreTeX' logo.

Mini-Profiler Plugin

- Built on top of Profiler Plugin
- Mimics StackExchange MiniProfiler (.NET and Ruby)
- Add `<miniprofiler:javascript/>` to `main.gsp`
- Heads up display
- Timing and SQL statements per artefact type



Mini-Profiler Plugin

- View SQL queries with actual params
 - great debugging tool
- No summary of repeated queries
- JS dialog a little flaky
- Still, recommended

Controller

T+259.0 ms

None

0.0 ms

```
select
  count(*) as y0_
from
  story this_
where
  this_.state_id=11
```

Controller

T+261.0 ms

None

0.0 ms

```
select
  this_.id as id2_0_,
  this_.version as version2_0_,
  this_.category as category2_0_,
  this_.name as name2_0_,
  this_.sort_order as sort5_2_0_
from
  board_state this_
where
  this_.category='ITERATION'
order by
  this_.sort_order asc
```

4.00 ms

Controller — 4.00 ms

Controller

T+265.0 ms

None

0.0 ms

```
select
  count(*) as y0_
from
  defect this_
where
  this_.state_id in (
    6, 7, 2, 12, 3, 8, 9, 4
  )
```

Hibernate Metrics

- ❑ Not yet released
- ❑ Focus is on Hibernate behavior and domain objects as well as SQL generated/executed
- ❑ Programmatically enables Statistics API, doesn't wrap DB driver
- ❑ Thus, able to turn on/off at will
- ❑ Heads up display ala Mini Profiler Plugin

Hibernate Metrics

- Timing and database/Hibernate info
- Intercepts logSql console output, groups queries by execution count
- Hibernate Stats doesn't report on criteria queries (until version 4.3+)

Before/After AJAX calls [Clear Metrics](#) [Refresh Metrics](#)

Time Metrics [ms](#) Total = **3215** Controller/Service = 86 View = **3129**

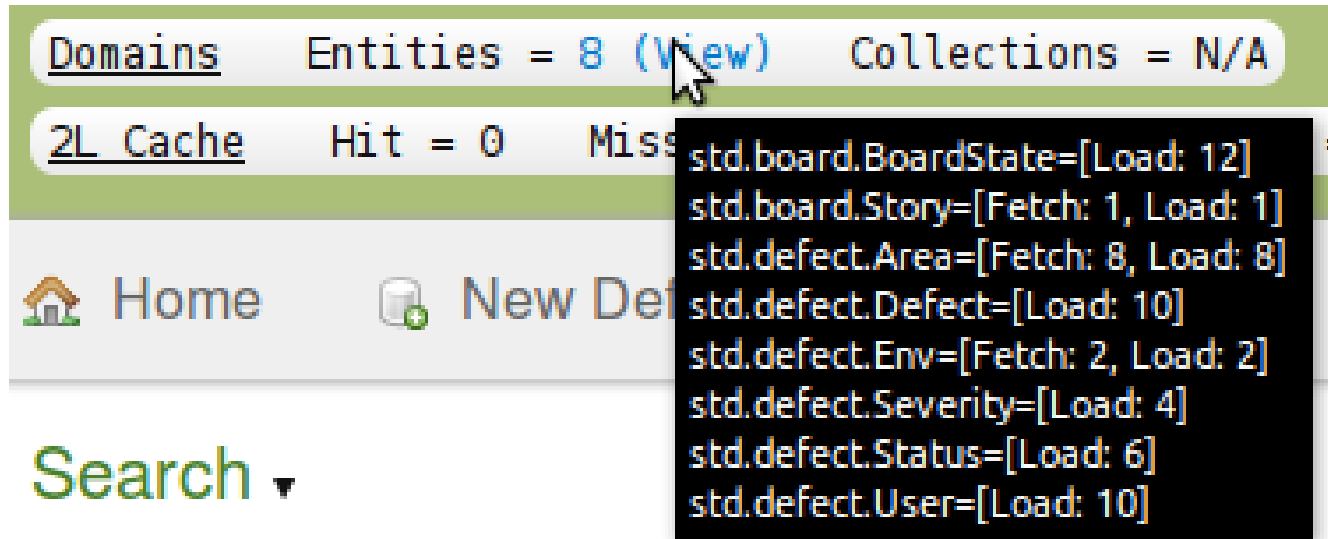
DB Metrics [Counts](#) Queries Executed = **66** Prepared Statements = **248** Transactions = 1 Flushes = 1

[SQL](#) Logged to Console = 20 ([View](#)) Executed = N/A Stats = N/A Slowest = N/A

[Domains](#) Entities = 13 ([View](#)) Collections = 5 ([View](#)) [Query Cache](#) Hit = 0 Miss = 0 Put = 0

[2L Cache](#) Hit = 0 Miss = 5 Put = 5 Domains = 2 ([View](#)) [Sessions](#) Opened = 2 Closed = 2

Hibernate Metrics



The screenshot displays the Hibernate Metrics web application. At the top, a green bar contains the text "Domains Entities = 8 (View) Collections = N/A". Below this, a white bar shows "2L Cache Hit = 0 Miss =". A mouse cursor is hovering over the "(View)" link, which has triggered a tooltip. The tooltip, with a black background and white text, lists the following metrics: std.board.BoardState=[Load: 12], std.board.Story=[Fetch: 1, Load: 1], std.defect.Area=[Fetch: 8, Load: 8], std.defect.Defect=[Load: 10], std.defect.Env=[Fetch: 2, Load: 2], std.defect.Severity=[Load: 4], std.defect.Status=[Load: 6], and std.defect.User=[Load: 10]. At the bottom left, there is a "Home" link with a house icon and a "New Defect" link with a document icon. A "Search" label with a dropdown arrow is also visible.

Domains Entities = 8 (View) Collections = N/A

2L Cache Hit = 0 Miss =

Home New Defect

Search ▼

std.board.BoardState=[Load: 12]
std.board.Story=[Fetch: 1, Load: 1]
std.defect.Area=[Fetch: 8, Load: 8]
std.defect.Defect=[Load: 10]
std.defect.Env=[Fetch: 2, Load: 2]
std.defect.Severity=[Load: 4]
std.defect.Status=[Load: 6]
std.defect.User=[Load: 10]

Hibernate Metrics

The screenshot displays a web application interface with a sidebar on the left and a main content area. The sidebar contains a 'DB Metrics' section with tabs for 'Counts', 'SQL', 'Domains', and '2L Cache'. The 'Counts' tab is active, showing 'Entities = 1' and 'Hit = 0'. Below this is a 'Severity List' section with a table containing two rows: 'Critical' and 'High'. The main content area shows a 'Logged to Console' dialog box with two SQL queries. The first query is a select statement with aliases, and the second is a count query. Both queries have an execution count of 1.

DB Metrics Counts SQL Logged to Console

Domains Entities = 1

2L Cache Hit = 0

Home New S

Severity List

Name
<u>Critical</u>
<u>High</u>

Logged to Console X

```
Hibernate: select
  this_.id as id15_0_,
  this_.version as version15_0_,
  this_.name as name15_0_
from
  severity this_ limit ?
[Execution Count: 1]
```

```
Hibernate: select
  count(*) as y0_
from
  severity this_
[Execution Count: 1]
```

Close

Hibernate Metrics – Integration

- Not actually released yet, so...
- Outside plugins block of BuildConfig

```
// BuildConfig.groovy
```

```
grails.plugin.location.'hibernate-metrics' =  
"/path/to/cloned/github/repo"
```

Hibernate Metrics – Integration

- Add display option to layout
- Turn on in dev via link
- Turn on in prod via URL

```
<!-- show normal header if not enabled -->  
<perfMetrics:isNotEnabled>  
    <a href="/"></a>  
</perfMetrics:isNotEnabled>
```

```
<!-- otherwise hide header, show stats -->  
<perfMetrics:isEnabled>  
    <perfMetrics:metrics />  
</perfMetrics:isEnabled>
```

```
<!-- enable/disable links in dev env -->  
<perfMetrics:devEnvControl />
```

```
// enable/disable directly (prod)  
/myApp/hibernateMetrics/enable  
/myApp/hibernateMetrics/disable
```


Hibernate Metrics – Integration

- Programmatic integration possible
- Primarily used from the Console Plugin

```
HibernateMetrics.withSqlLogging {  
    // execute code that you want SQL to be logged for  
}  
  
// println output  
Time Metrics:  
    Total Time (ms) = 30  
    Controller/Service (ms) = 28  
    View (ms) = 2  
DB Metrics:  
    Total Queries = 2  
    Prepared Statements = 2  
    Logged SQL = ...  
    Entity Info = std.board.BoardState = [Load: 12]  
    ...  
    Sessions Opened = 0  
    Sessions Closed = 0  
    Transaction Count = 0  
    Flush Count = 0
```




Demo

Agenda

- The landscape
- The solutions
- The problems that the solutions have uncovered
- And some fixes

Some of the Things I've Learned




I AM GORM

GREAT AND POWERFUL


Environments

80+ domains



Environments

190+ domains



Queries in a Loop

- Easy win to minimize number of round trips to the DB

```
// results in query per id  
params.inputIds.each {  
  list << Defect.get(it)  
}
```

```
// use .getAll() instead - single query  
list << Defect.getAll(params['inputIds'])
```

Queries in a Loop - Continued

- Pay attention any time you're looping
- Queries may be wrapped/obfuscated by other logic

```
// this logic works, but can we be smarter about it?
```

```
params.inputIds.each { id ->  
  story.addToDefects(  
    buildDefectWithSpecialLogic(  
      Defect.get(id), user, client))  
}
```

```
// single query instead
```

```
List d = Defect.getAll(params.inputIds)  
d?.each { defect ->  
  story.addToDefects(  
    buildDefectWithSpecialLogic(  
      defect, user, client))  
}  
}
```


Navigating Obj Structure To Get Minimal Data

- A lot of data is a terrible thing to waste

```
// loads full Story, Status, Comment objects  
Story.findAllByStatus(  
  Status.findByName("Completed")  
)*.comments*.user?.name
```

```
// loads just the data needed, in one query  
Story.createCriteria.list() {  
  comments {  
    user {  
      property('name')  
    }  
  }  
  status {  
    eq('name', 'Completed')  
  }  
}
```

Navigating Obj Structure - Continued

- Happens in other scenarios too

```
// a similar thing
story.subTasks*.estimates.find {
  it.id == id
}

// loads just the data needed, in one query
Story.createCriteria.get() {
  subTasks {
    estimates {
      eq('id', id)
    }
  }
}
```

hasOne Always Eager Fetches

- If the OtherDomain in the hasOne is expensive to load this can be an unwelcome surprise

```
class MyDomain {  
  static hasOne = [alwaysAlongForTheRide:  
    OtherDomain]  
}  
  
class OtherDomain {  
  static hasOne = [a:A, b:B, c:C, d:D, ...]  
  static mapping = {  
    collection1 fetch:'join'  
    collection2 fetch:'join'  
    collection3 fetch:'join'  
    collection4 fetch:'join'  
    collection5 fetch:'join'  
    ...  
  }  
}
```

N+1 Queries Problem

- Default per relationship is lazy loading

```
class Story {  
  static hasMany = [defects:Defect]  
}  
  
// iterating over collection causes N+1  
Story.get(id)?.defects?.each { defect →  
  hoursCount += defect.hoursSpent  
}
```

Fixing N+1 Queries

- ❑ Fetch:'join' – join table & get all data in a single query
- ❑ Lazy:false – immediately execute second query to load collection
- ❑ Override per query


```
class Story {  
  static hasMany = [defects:Defect]  
  
  // pick one  
  static mapping = {  
    //defects lazy:false  
    //defects fetch:'join'  
  }  
}  
  
// specify at query time  
Story.list( [fetch:[defects:'join']] )
```

Another N+1 Gotcha (?)

- Not naming `belongsTo` seems to guarantee N+1 queries

```
class Story {  
  static hasMany = [defects:Defect]  
}  
  
// creates a join table  
class Defect {  
  static belongsTo = [Story]  
}  
  
// no join table, story_id on Defect table  
class Defect {  
  static belongsTo = [story:Story]  
}
```

Real World Example



Custom Validators

- Validators can get called a lot...
- At the very least use `save(validate:false)` if manually checking `validate()`

```
class Story {  
  static constraints = {  
    name validator: { val, obj ->  
      doQueryToCheckUniqueness(val, obj)  
    }  
  }  
}  
  
// save() will call validate by default  
if (story.validate()) {  
  // so don't call it twice  
  story.save(validate:false)  
}
```


Where Queries Don't Use 2nd Level Cache

- Where queries create DetachedCriteriaObject
- Not associated to a Hibernate Session
- Thus, no second-level cache

```
class MyConfig {  
    static mapping = {  
        cache true // use 2nd level cache  
    }  
}  
  
// nice syntax, always going to hit the DB  
def configVal = MyConfig.where {  
    name == configName  
    customer == currentCustomer  
    user == user  
}.get()
```

Another 2nd Level Cache Issue

- A misplaced colon makes a big difference

```
class MyConfig {  
    static mapping = {  
        cache: true // fail  
    }  
}
```

Loading Data, Ignoring Data

- View gets refactored
- Controller or service is still retrieving data

```
// controller action
def show() {
  [users:User.list(),
   severities:Severity.list(),
   environments:Env.list()]
}

// view
<g:each var="user" in="${users}">
  ..
</g:each>

<g:each var="severity" in="${severities}">
  ..
</g:each>
```

Collections load all instances when adding

- hasMany are Sets by default
- Set guarantees uniqueness
- Really only comes into play with huge collections

When Groovy SQL Is Better

- Batch jobs
 - especially big, interrelated data
- Real world example
 - converting from domains, dynamic finders, etc to Groovy SQL
 - job went from 1 week+ to 2-4 hours
- Have to do some things manually
- Pay attention to the downsides

```
// possible SQL injection sql.execute('select  
* from x where y = ' + input)
```

```
// avoids it  
sql.execute('select * from x where y = ? and z  
= ?', [inputA, inputB])
```

```
// pay attention to String vs GString  
sql.eachRow("select * from x where y = ${y}")  
// won't work
```

```
sql.eachRow("select * from x where y = ${y}"  
as String)  
// will work
```

Fix It, Yo

- Don't prematurely optimize
- Proper DB indexes for your queries
- Warm it up – or cool it down
- A lot of it is “it depends”

Hibernate Metrics – The Future

- Data is not tracked over time
- Charts, graphs
- Integration with Integration Tests
- Actually release the damn thing



Thank You For Not Sleeping

(I hope)

Questions?

Links / Credits

- Hibernate Metrics Plugin
<https://github.com/davidkuster/hibernate-metrics>
- Ted Neward – The Vietnam of Computer Science
<http://blogs.tedneward.com/2006/06/26/The+Vietnam+Of+Computer+Science.aspx>
- Burt Beckwith – Advanced GORM – Performance, Customization and Monitoring
<http://infoq.com/presentations/GORM-performance>
- StackExchange MiniProfiler Plugin
<http://miniprofiler.com>
- Tom Dunstan – Debugging Grails Database Performance
<http://skillsmatter.com/skillscasts/3785-debugging-grails-database-performance>
- Slides
<http://talldave.net>

Links / Credits

- The Wonderful Wizard of Oz
<http://www.read.gov/books/oz.html>
- The Wizard of Oz
<http://thewizardofoz.warnerbros.com>
- Flying Spaghetti Monster
<http://dextermurphy.deviantart.com/art/Flying-Spaghetti-Monster-244993860>