

# Why your build matters

Peter Ledbrook

e: [peter@cacoethes.co.uk](mailto:peter@cacoethes.co.uk)  
w: <http://www.cacoethes.co.uk>  
t: @pledbrook

# When did it all begin?



- Programmable machines have a long history
- Mechanised instruments
- Looms



## Punched cards

First used for looms in 18th Century France – Jacquard Loom 1801

IBM introduced punched card format 1928 – 80 columns!

Text mode column width for DOS, Unix etc.

Imagine creating by hand

Keypunch machines (a little like typewriters) for creating them

Errors mean throwing away the card and doing it again

Fortran and Cobol started with punched cards

# Fixing card decks



What if your cards are dropped on the floor?  
This is an IBM 082 Sorter  
Automation of a very error-prone manual process

# Magnetism saves the day!



Finally, R/W data and code. The dawn of text editors.

# C/C++

Source



Compiler



Linker



.so/.dll & .exe

**Multi-platform  
builds!**

# Make

```
*.c: *.o
```

```
cc ...
```

```
myapp: app.o mylib.o ...
```

```
ld ...
```

+ manual custom dependencies

Focus mainly on compilation and linking

Make has powerful model based on file dependencies

Maintenance was a lot of work

# Make

```
*.c: *.o
```

```
cc ...
```

```
myapp: app.o mylib.o ...
```

```
ld ...
```

What's this?


+ manual custom dependencies



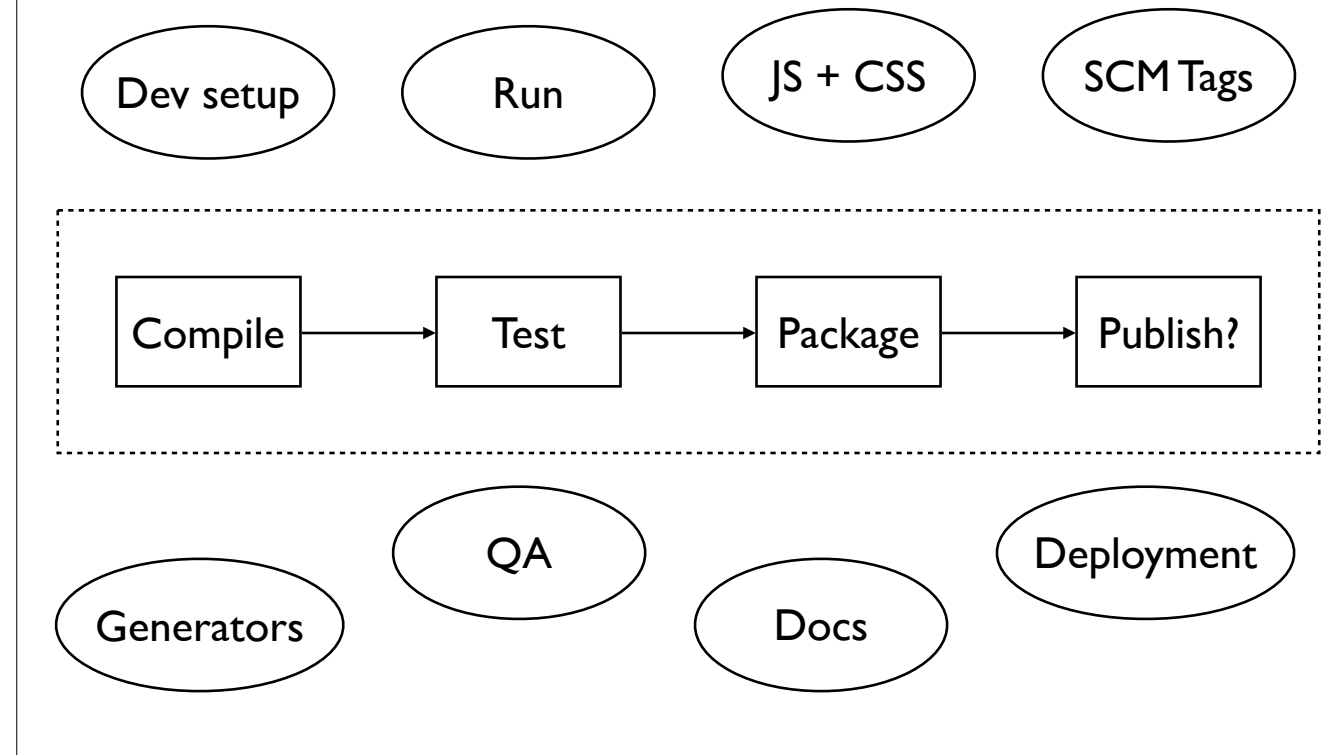
# Java

- Compiler handles .java to .class dependencies
- JAR as 'module' unit
- Several standard deployment types
  - Applet
  - WAR
  - Standalone app

# A standard build




# Java build evolved



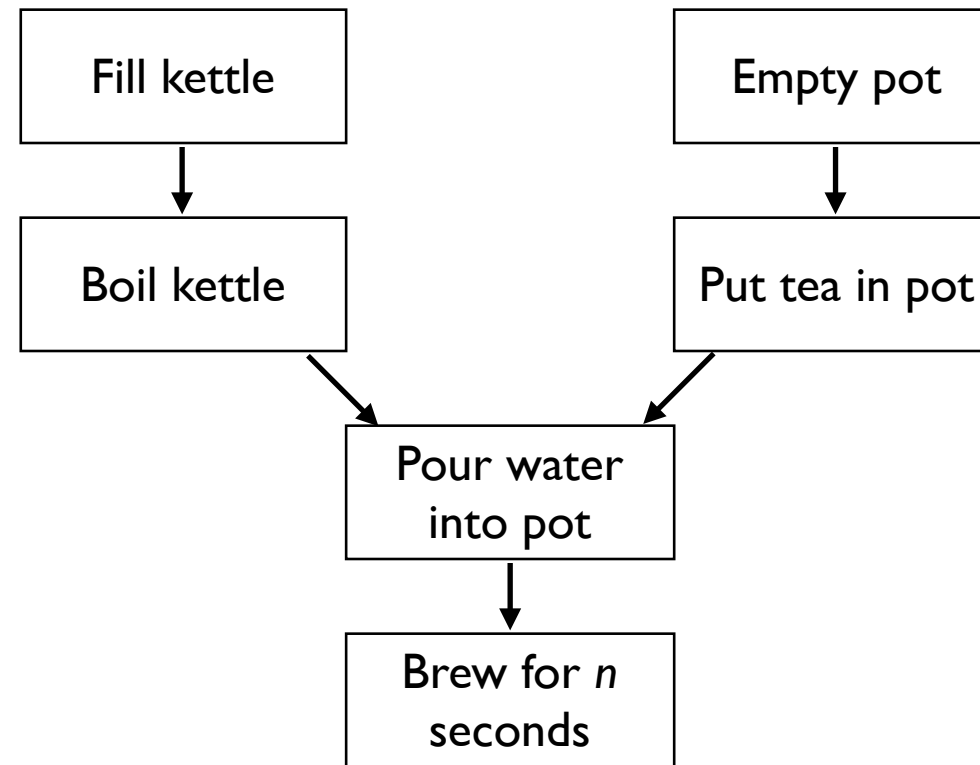
Builds are incorporating more and more steps

**What will builds look like in 5  
years time?**




What is common across all these systems?  
To understand that, let's look at a non-software process

# Making tea



Clear that it's a series of steps (or tasks)  
Some tasks require others to complete first  
Others can run in parallel  
Acyclic graph of tasks!

# Making tea (my way)

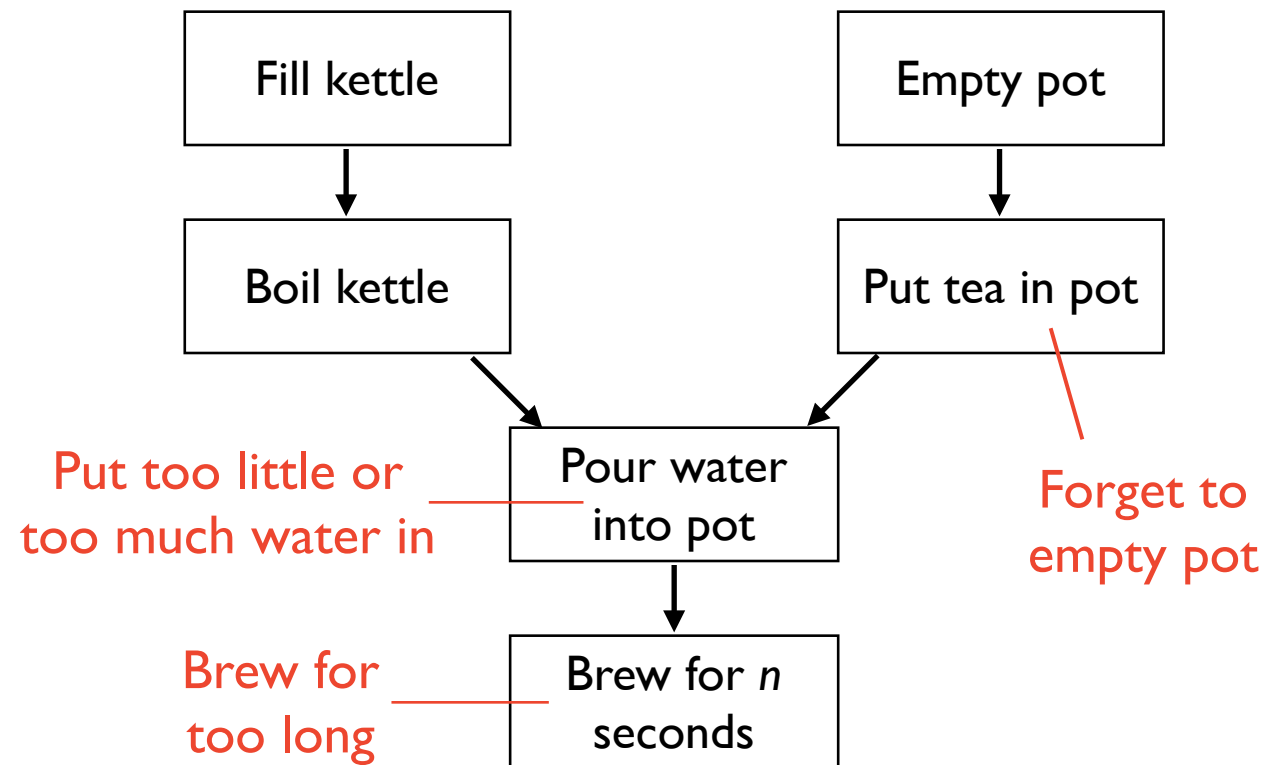


A different way to make tea  
Less standard, more appropriate for me

**A standard process does not  
mean an exclusive one**



# Many things can go wrong



# Automate to eliminate human error

People are fallible

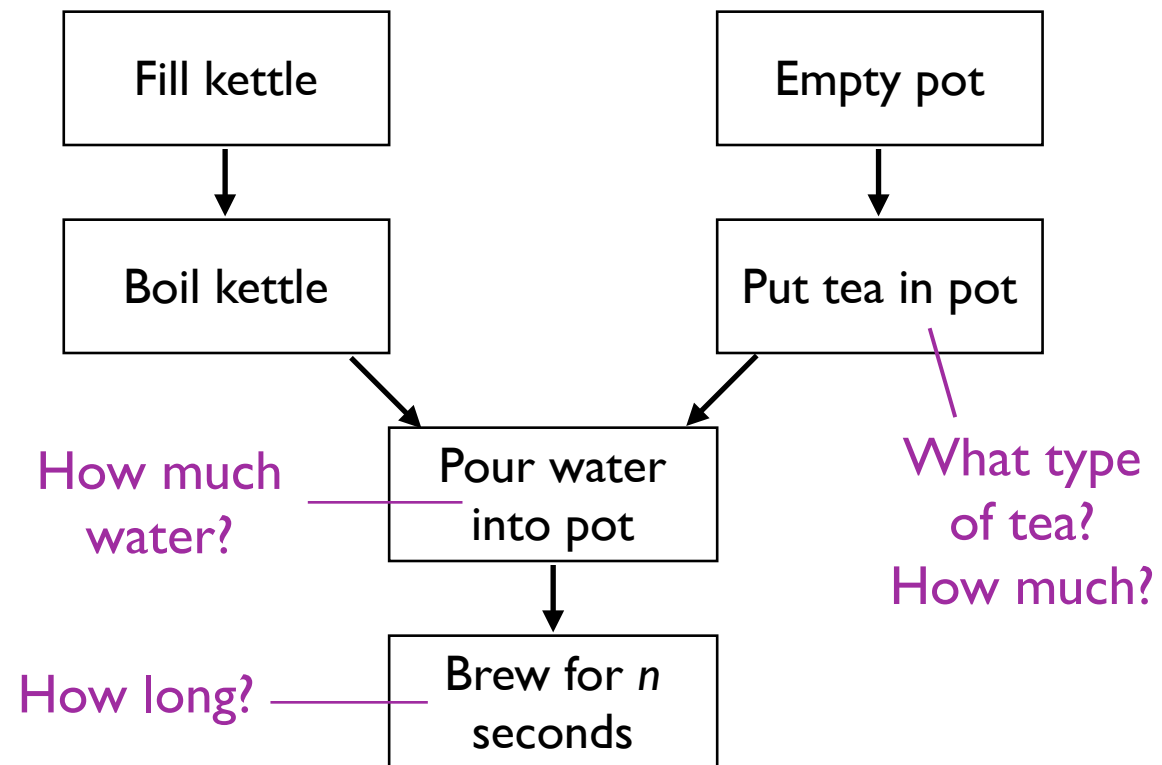
Muscle memory can help

e.g. tying shoe laces

Not usually applicable in software development

What if your job depends on  
perfect tea every time?

# Inputs & parameters



Build inputs:

- \* water quality
- \* tea bags/loose leaf
- \* tea type

Build parameters:

- \* quantity of tea
- \* quantity of water
- \* brew time

Same inputs should result in  
same outputs

In other words, repeatable builds

Environment should not be a factor

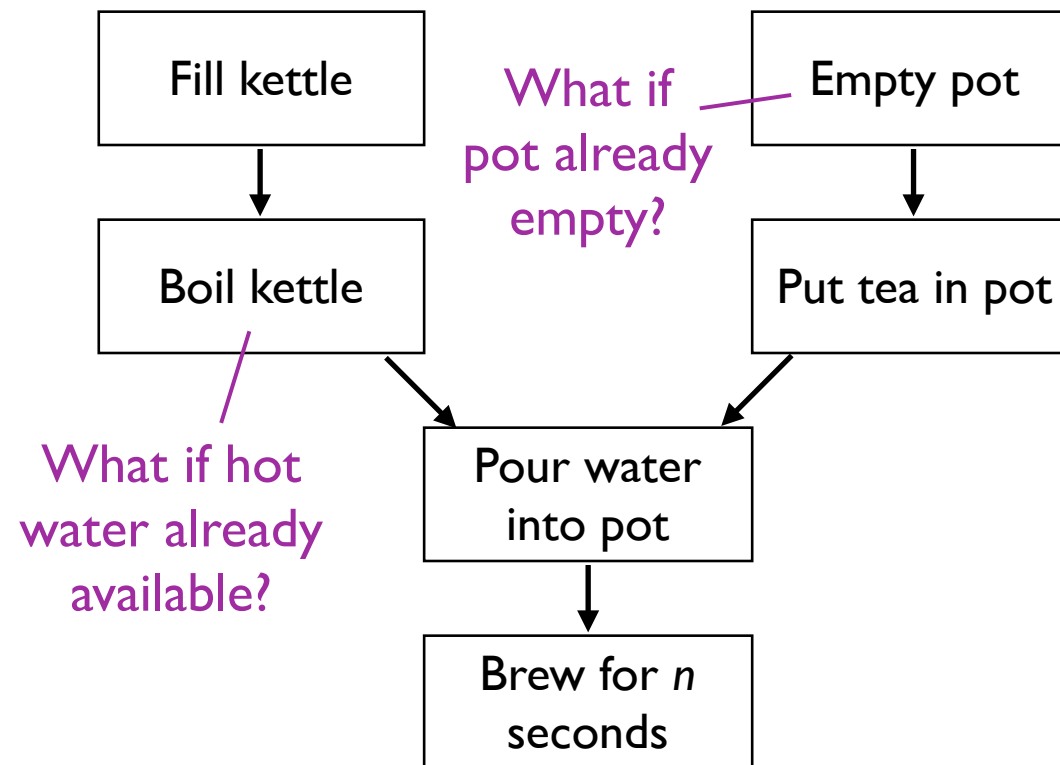
- boiling kettle at top of mountain reduces temperature of water

Consider “works for me” with local Maven cache

How many cups of tea do you  
need to make a day?

Speed and efficiency are useful

# Task avoidance



Doing all steps slows you down if some are unnecessary  
Output of “boil kettle” is hot water  
If hot water already exists, no need to boil

**Incremental build (up-to-date  
checks) saves time**

Often quicker to check whether a task is up to date rather than run it regardless



A build should be...

Automated

Repeatable

As fast as possible

The

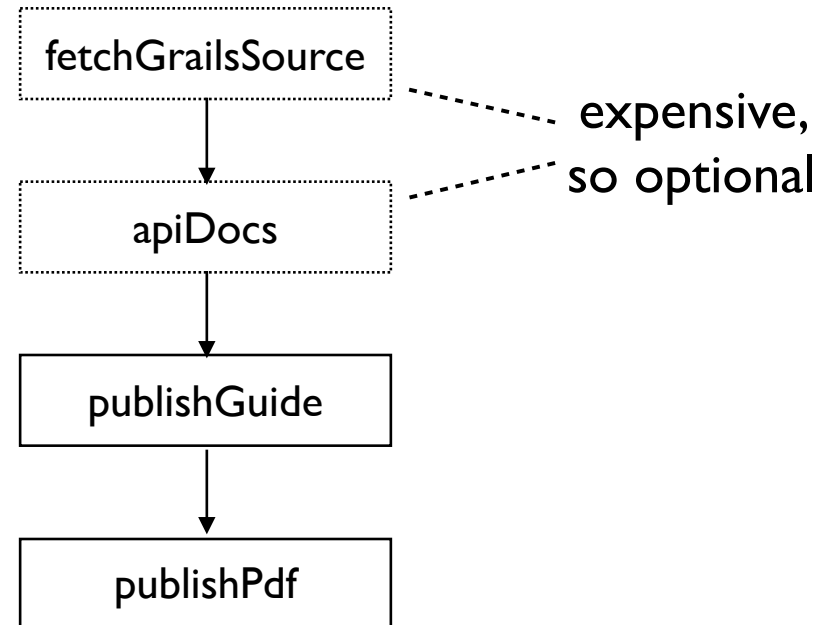


Way

# Automation

# Custom build

grails-doc  
project



A build to generate the Grails user guide  
Automation of all steps  
No standard tasks  
Task graph is a suitable model

# Custom tasks in Gradle

```
task publishGuide << {  
    // Generate HTML from .gdoc  
}
```



buildSrc/src/groovy/pkg/PublishGuideTask.groovy



pkg.PublishGuideTask in a JAR

# Optional task

```
apiDocs.onlyIf {  
    !System.getProperty("disable.groovydocs")  
}
```

# Add custom tasks to standard Groovy/Java projects

Even “standard” builds often have custom steps, e.g.

- integration tests

- deployment

- documentation

Stop using scripts or other tools separate from the build!

# Repeatable builds



# Depends on tasks

Are tasks environment-dependent?

Do system properties or environment variables have an effect?

What about number of cores?

Order of tests?

Gradle can't help much with this – except for task ordering

# Task ordering

- mustRunAfter/shouldRunAfter
- No task dependencies

```
task integTest {  
    ...  
}  
  
task packageReports {  
    mustRunAfter integTest  
    ...  
}
```

`packageReports` does not depend on `integTest`  
But if both are executed, `packageReports` must come after `integTest`  
`shouldRunAfter` is less strict – mostly to optimise feedback

# Task ordering

- finalizedBy
- Ensures execution of the finalizer task

```
task integTest {  
    finalizedBy packageReports  
    ...  
}  
  
task packageReports {  
    ...  
}
```

`packageReports` does not depend on `integTest`  
If `integTest` runs, then `packageReports` will run too, always after `integTest`

# Dependency issues

- Maven cache pollution (“works for me” syndrome)
- Different remote repository configurations (where is the dependency coming from?)
- Version resolution, eviction, and failed eviction
  - Multiple JARs on the classpath and order counts

# Gradle cache

- Origin checks
- Artifact checksums
- Concurrency safe
- Avoid `mavenLocal()`!

Artifacts are stored with source repo URL (allows origin checks)

Artifact checksums protect against different binaries with same name & version

Doesn't solve the version conflicts issue

# Resolution strategy

Defaults to 'newest' strategy

```
configurations.all {  
    resolutionStrategy {  
        failOnVersionConflict()  
  
        force 'asm:asm-all:3.3.1',  
              'commons-io:commons-io:1.4'  
    }  
}
```

Fail the build if any  
version conflicts

Override version to use  
for specific dependencies

Fine-grained control over dependency versions  
Automatic conflict resolution error-prone

# Resolution strategy

```
configurations.all {  
    eachDependency { details ->  
        if (details.requested.name == 'groovy-all') {  
            details.useTarget(  
                group: details.requested.group,  
                name: 'groovy',  
                version: details.requested.version)  
            }  
        }  
    }  
}
```

\  
Control modules with  
different names, same classes

A painful problem to debug

# Debugging dependencies

gradle dependencies

gradle dependencyInsight --dependency ...

```
configurations.all.files.each { File f ->
    println f.name
}
```

`dependencies` gives you all dependencies in your project

`dependencyInsight` shows why/how a given dependency is included in the project

`<conf>.files` lists all the files and directories that will be on configuration's classpath



**Fast build execution**

# Depends on tasks

How long do individual tasks take?  
No parallel execution of tasks currently  
Can build decoupled projects in parallel

# Incremental build

```
:lazybones-app:compileJava UP-TO-DATE  
:lazybones-app:compileGroovy UP-TO-DATE  
:lazybones-app:processResources UP-TO-DATE  
:lazybones-app:classes UP-TO-DATE  
:lazybones-app:jar UP-TO-DATE  
:lazybones-app:startScripts UP-TO-DATE  
:lazybones-app:installApp UP-TO-DATE
```

BUILD SUCCESSFUL

Total time: 2.178 secs

Don't execute tasks you don't have to

# Task inputs & outputs

```
class BintrayGenericUpload extends DefaultTask {  
    @InputFile File artifactFile  
  
    @Input      String artifactUrlPath  
  
    @Optional  
    @Input      String repositoryUrl  
    ...  
}
```

Use of annotations makes task support incremental build  
Inputs and outputs can be values, files, directories

**Who** needs to use the build?

**What** do they need to use it?

Will they use the **whole build**?

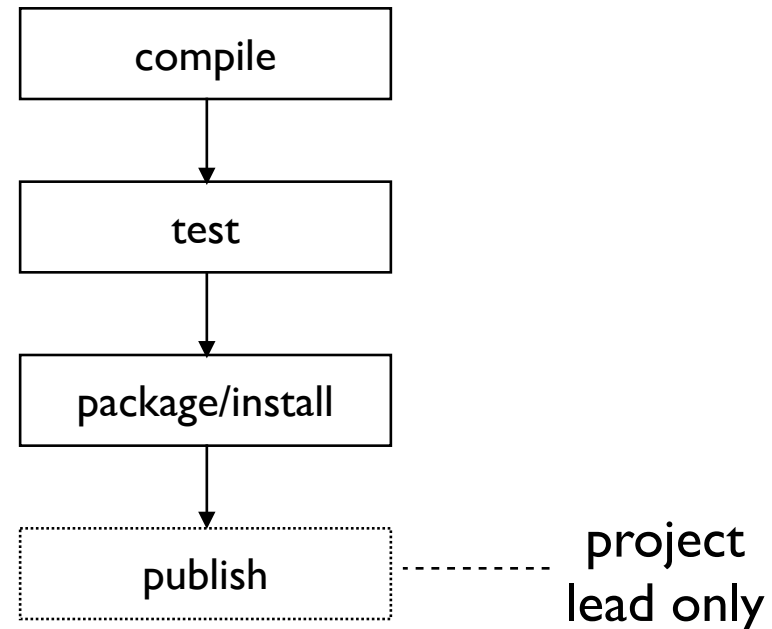
**How long** does it take?

Developers, QA, production types

How much software needs to be set up before someone can use the build? How long is your wiki page?

# User categories

Lazybones  
project



If the project is published, require credentials  
Fail fast  
Other users don't require credentials

# Interrogate task graph

```
gradle.taskGraph.whenReady { graph ->
    if (graph.hasTask(":lazybones-app:uploadDist")) {
        verifyProperty(project, 'repo.url')
        verifyProperty(project, 'repo.username')
        verifyProperty(project, 'repo.apiKey')

        uploadDist.repositoryUrl = project.'repo.url'
        uploadDist.username = project.'repo.username'
        uploadDist.apiKey = project.'repo.apiKey'
    }
}
```

Only fail fast if `uploadTask` will be executed

Fail slow would require integration tests to run – adding minutes to deployment

# Other uses

- Include class obfuscation conditionally
- Limit visibility of tasks based on role
- Update version based on 'release' task



**Finally...**

# A build should have...

- Features
- Error reporting
- A model of the process

Just like any other piece of  
software!

# Summary

- End-to-end process is the build
- 80-20 rule (80% standard 20% custom)
- Automate everything == save time
- Invest in build as if it's part of your code base
- Building software requires a rich model